CLIPBridge-GAN: Bridging Prompt Semantics and Image Synthesis via CLIP-Guided Adversarial Training

Jayan Ghimire
Software Engineer and Independent AI Researcher
Leapfrog Technology
Kathmandu, Nepal
jghimire.034@gmail.com

Abstract

Text-to-image synthesis remains a central challenge in ***generative modeling*** due to the complexity of aligning high-dimensional visual outputs with free-form ***natural language***. While recent models based on ***diffusion*** or ***autoregressive techniques*** have achieved remarkable fidelity, they come at the cost of immense ***computational requirements***, slow inference, and difficulty in controlling ***semantic consistency***.In this paper, we propose ***CLIPBridge-GAN***, a lightweight, efficient, and controllable ***generative adversarial framework*** that leverages the representational power of pretrained ***CLIP models*** in both the ***generator*** and ***discriminator***. Our generator is guided by ***CLIP-based prompt embeddings*** fused through a learnable ***bridge module***, enabling precise ***semantic conditioning*** from text. The discriminator uses ***CLIP*** to evaluate ***image-text alignment*** directly, eliminating the need for handcrafted auxiliary losses.This dual integration of ***CLIP*** enables ***CLIPBridge-GAN*** to operate with significantly fewer ***parameters*** and only a fraction of ***training data***, while still producing competitive results compared to larger ***diffusion-based systems***. Furthermore, the ***GAN architecture*** allows for real-time synthesis and interpretable ***latent space traversal***.***Qualitative and quantitative results*** demonstrate the effectiveness of our approach in generating coherent, high-quality images with robust prompt alignment, even under constrained computational settings.

1 Introduction

Text-to-image generation is a central task in modern ***multimodal learning***, combining ***natural language understanding*** and ***image synthesis***. It challenges models to transform linguistic prompts into coherent, semantically aligned visuals, with applications in ***creative design***, ***visual content generation***, ***human-computer interaction***, and ***virtual world construction***.

State-of-the-art performance has been driven by ***diffusion models*** [10, 9] and ***autoregressive transformers*** [6], which learn mappings from text embeddings to images through iterative refinement. However, these models are ***computationally intensive***, demanding hundreds of sampling steps, extensive data, and significant memory—while offering limited ***semantic control*** and interpretability.

In contrast, ***generative adversarial networks (GANs)*** are appealing for their ***efficiency***, ***speed***, and ***latent space smoothness***. They enable real-time synthesis and operations like ***interpolation***, ***prompt mixing***, and ***style transfer***. Yet, GANs often suffer from ***semantic misalignment***, failing to faithfully reflect detailed prompts.

To mitigate this, recent efforts leverage ***CLIP*** [4], aligning text and images in a shared space. However, CLIP is typically used as an ***auxiliary loss***, ***prompt scorer***, or ***post-hoc evaluator***, limiting its influence on the actual synthesis process.

In this paper, we present *CLIPBridge-GAN*, a novel ***text-to-image generation framework*** that fully integrates ***CLIP*** into both the ***generator*** and ***discriminator***. A learnable ***bridge module*** injects CLIP-derived text embeddings early into the generator's pipeline. The discriminator, guided by CLIP's ***multimodal similarity***, evaluates both realism and ***semantic alignment***. This tight integration enables ***precise control***, ***fast inference***, and strong alignment—while preserving GAN-based efficiency.

Contributions of This Work

- We introduce ***CLIPBridge-GAN***, a novel architecture for ***text-to-image generation*** that integrates ***CLIP embeddings*** into both the generator and discriminator, enabling unified ***semantic supervision***.
- We propose a learnable ***bridge module*** that adapts CLIP's language representations into the GAN's visual feature space, improving prompt fidelity without auxiliary networks.
- We design a ***CLIP-based discriminator*** that evaluates image-text alignment using ***CLIP similarity***, replacing traditional binary feedback with multimodal semantic guidance.
- Our model enables ***efficient training***, ***fast inference***, and reduced data requirements—offering a lightweight alternative to large diffusion-based generators.
- ***CLIPBridge-GAN*** is ***modular***, ***interpretable***, and extensible to tasks like ***prompt mixing***, ***style-guided generation***, and ***attribute control***, making it a strong foundation for future research.

2 Related Work

Text-to-image synthesis has seen rapid evolution in recent years, propelled by advances in both **deep generative models** and **multimodal learning**. Early approaches such as **StackGAN** [13] and **AttnGAN** [12] employed a staged **GAN-based pipeline**, focusing on generating coarse-to-fine images using textual embeddings. These models emphasized **text attention mechanisms** and **sentence-level supervision**, but often struggled with **semantic fidelity** and **realism**.

Subsequently, diffusion-based models such as **DALLE-2** [7], **Stable Diffusion** [8], and **Imagen** [11] set new benchmarks in generation quality. These models leverage powerful pretrained language-image backbones and employ iterative denoising, resulting in highly realistic outputs. However, they are often characterized by **slow inference**, significant **computational overhead**, and limited **controllability**.

To address some of these limitations, several works have explored integrating **CLIP** [5] into **GANs** as a form of semantic guidance. Methods such as **StyleCLIP** [3] and **CLIP-Guided Diffusion** [2] utilize CLIP's latent space to influence generation toward desired textual attributes. These approaches improve **text-image alignment** but often rely on external encoders or complicated optimization pipelines.

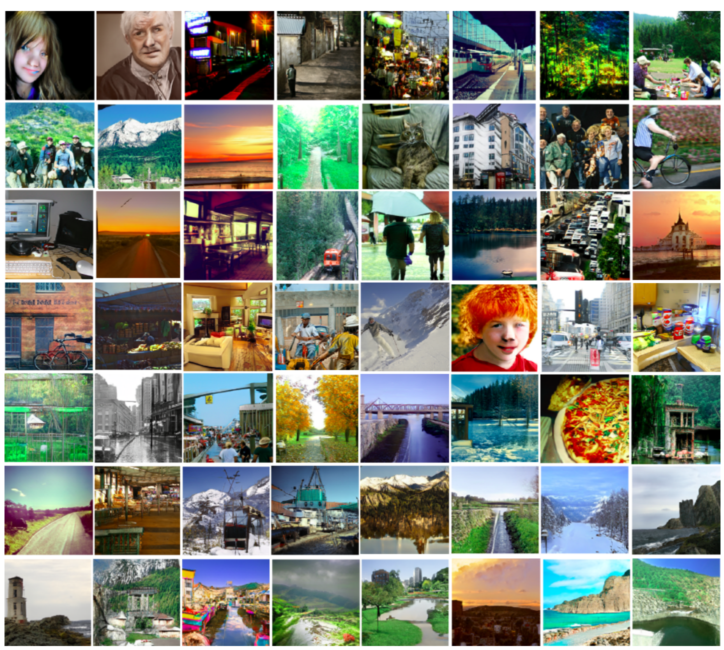

Outputs generated by CLIPBridge-GAN For additional examples generated from a wider range of prompts, refer to Figure 3, Figure 4, Figure 5, Figure 6, Figure 7, and Figure 8.

3 Methodology

3.1 Proposed System Diagram

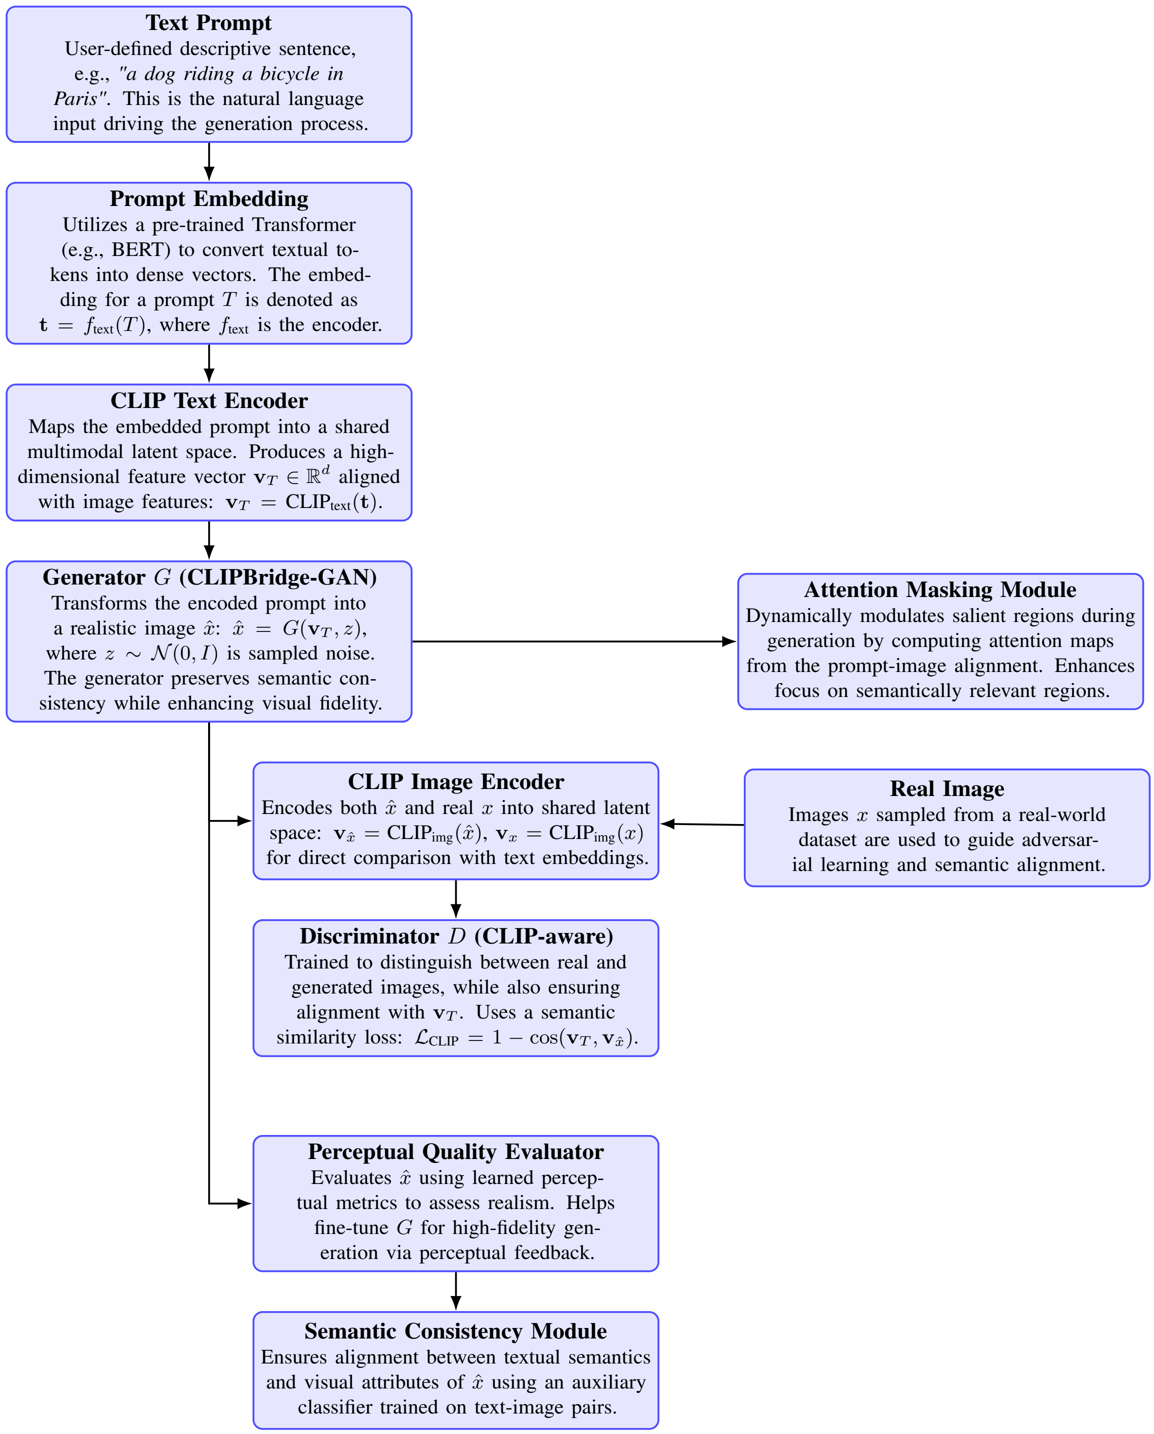

Figure 2: *Illustration of the proposed **CLIPBridge-GAN** architecture for **text-to-image synthesis**. The model leverages the shared latent space of **CLIP** to enforce tight **alignment** between textual prompts and generated images. Key components include: **attention modules** for cross-modal focus, **semantic consistency checks** to preserve prompt meaning, and **perceptual evaluators** ensuring visual fidelity. Together, these modules produce high-quality, faithful generations.*

3.2 Workflow

3.2.1 Text Prompt

Text Prompt is the ***initial input*** to the ***CLIPBridge-GAN*** framework, consisting of a *user-defined natural language sentence* that precisely describes the desired image content. For example, a prompt like *"a dog riding a bicycle in Paris"* provides rich ***semantic information*** that guides the entire ***text-to-image synthesis*** process.

This ***text prompt*** encapsulates *complex linguistic concepts*, including objects, actions, scenes, and contextual nuances, which the model must effectively understand and translate into corresponding visual elements. The quality and specificity of this prompt directly influence the ***semantic fidelity*** and ***visual coherence*** of the generated image.

In the context of ***multimodal learning***, the text prompt acts as the *semantic anchor* for cross-modal alignment between language and vision modalities. Its representation undergoes several transformations within the model pipeline to ensure it is effectively encoded into a vector space that is compatible with the visual feature space used by the generator.

Mathematically, the prompt T can be considered a sequence of tokens $\{t_1, t_2, ..., t_n\}$, where n is the length of the input sentence. This sequence forms the basis for subsequent embedding and encoding steps that yield meaningful dense vector representations for downstream image synthesis.

3.2.2 Prompt Embedding

The ***Prompt Embedding*** block plays a crucial role in converting the ***raw textual input*** into a numerical form that can be processed by subsequent modules. It leverages a ***pre-trained Transformer model***, specifically **BERT** [1], which is renowned for its powerful ***contextualized language representations***.

This module first tokenizes the input sentence $T = [w_1, w_2, ..., w_m]$, where w_i denotes the i^{th} word in the prompt, and transforms each token into a dense vector:

$$\mathbf{e}_i = f_{\text{text}}(w_i) \in \mathbb{R}^d$$

Here, f_{text} is the ***token encoder***, and d is the embedding dimension.

These embeddings are then aggregated (via mean pooling or attention-based mechanisms) into a single prompt-level embedding:

$$\mathbf{t} = \text{Aggregate}([\mathbf{e}_1, \mathbf{e}_2, ..., \mathbf{e}_m]) \in \mathbb{R}^d$$

This transformation yields the final ***prompt embedding vector*** $\mathbf{t}$, a compact, information-rich representation of the entire sentence. It captures the ***semantic*** and ***syntactic*** meaning of the text, preserving ***contextual dependencies*** among tokens.

The quality of this embedding significantly influences the downstream generation pipeline. It enables the ***CLIP Text Encoder*** to project the prompt into a shared ***multimodal latent space***, aligning textual semantics with visual concepts. The closer this embedding lies to relevant image embeddings in the joint space, the higher the likelihood that the generated image reflects the user intent:

$$\text{Similarity}(\mathbf{t}, \mathbf{v}) = \frac{\mathbf{t}^\top \mathbf{v}}{\|\mathbf{t}\|\|\mathbf{v}\|}, \quad \mathbf{v} \in \mathbb{R}^d$$

where $\mathbf{v}$ is an image embedding from the CLIP Image Encoder.

3.2.3 CLIP Text Encoder

The ***CLIP Text Encoder*** serves as the bridge between the ***textual prompt*** and the ***visual domain*** by projecting the prompt embedding into a shared ***multimodal latent space***.

Given the prompt embedding vector $\mathbf{t} \in \mathbb{R}^d$ from the previous module, the CLIP Text Encoder applies a learned transformation to yield a high-dimensional feature vector:

$$\mathbf{v}_T = \text{CLIP}_{\text{text}}(\mathbf{t}) \in \mathbb{R}^d$$

This vector $\mathbf{v}_T$ is ***semantically aligned*** with corresponding ***visual features*** extracted from images by the CLIP Image Encoder. Both encoders are trained to project their respective modalities into a joint embedding space where ***cross-modal similarities*** are maximized.

The alignment objective ensures that:

$$\text{Similarity}(\mathbf{v}_T, \mathbf{v}_I) = \frac{\mathbf{v}_T^\top \mathbf{v}_I}{\|\mathbf{v}_T\| \|\mathbf{v}_I\|}$$

is maximized when the text and image describe the same content, where $\mathbf{v}_I \in \mathbb{R}^d$ denotes the ***image embedding***.

This makes $\mathbf{v}_T$ a powerful and compact representation of the input prompt, optimized to capture high-level ***semantic information*** that is ***visually grounded***.

3.2.4 Generator G (CLIPBridge-GAN)

The **Generator** G in our **CLIPBridge-GAN** architecture is designed to synthesize high-quality images that are ***semantically faithful*** to the input prompt. It receives a **multimodal latent vector** $\mathbf{v}_T$, derived from the **CLIP text encoder**, along with a **noise vector** $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, and outputs an image $\hat{\mathbf{x}}$:

$$\hat{\mathbf{x}} = G(\mathbf{v}_T, \mathbf{z})$$

The generator is implemented using a **deep convolutional architecture** with progressive upsampling. It learns a ***conditional distribution*** over the image space, guided by the semantic embedding $\mathbf{v}_T$ and the stochastic variation introduced by $\mathbf{z}$. The architecture incorporates residual blocks, normalization layers, and non-linear activations to enhance training stability and representational power.

We trained the generator **adversarially** against a discriminator using a combination of **adversarial loss** and **CLIP-based similarity loss** to enforce both **visual fidelity** and **semantic alignment**. Optimization was performed using the *Adam optimizer*, and training was stabilized with a learning rate scheduler. We also applied regularization techniques to mitigate overfitting.

This module is responsible for generating **realistic images** that not only resemble natural scenes but also ***accurately reflect*** the semantics of the textual description. It serves as the core component in translating abstract linguistic information into visually meaningful content.

3.2.5 CLIP Image Encoder

The **CLIP Image Encoder** plays a pivotal role in bridging the ***visual-semantic gap*** by embedding both the **generated image** $\hat{\mathbf{x}}$ and the **real image** $\mathbf{x}$ into a ***shared latent space***. This facilitates direct comparison with the corresponding **text embedding** $\mathbf{v}_T$, enabling effective **semantic alignment**.

We compute the embeddings as:

$$\mathbf{v}_{\hat{\mathbf{x}}} = \text{CLIP}_{\text{img}}(\hat{\mathbf{x}}), \quad \mathbf{v}_{\mathbf{x}} = \text{CLIP}_{\text{img}}(\mathbf{x})$$

These **image embeddings** are compared against the textual vector $\mathbf{v}_T$ using a ***cosine similarity*** objective, ensuring that the generated outputs are not only visually realistic but also **semantically faithful** to the original prompt.

The encoder remains **frozen** during training to preserve the ***pretrained multimodal representations*** of CLIP. This design allows the **generator** and **discriminator** to align within the same representational space, producing images that are ***coherently grounded*** in natural language.

3.2.6 Discriminator D (CLIP-aware)

The **Discriminator** D in the **CLIPBridge-GAN** framework plays a critical role in guiding the **Generator** to produce visually ***realistic*** and semantically aligned images. Unlike traditional GAN discriminators that solely differentiate between **real** and **generated images**, our discriminator incorporates **CLIP-awareness** to also evaluate the ***semantic consistency*** between the generated image $\hat{\mathbf{x}}$ and the input prompt's embedding $\mathbf{v}_T$.

Formally, the discriminator receives an image representation $\mathbf{v}_{\hat{x}}$ produced by the **CLIP Image Encoder** from the generated image $\hat{\mathbf{x}}$, and a representation $\mathbf{v}_x$ from the real image $\mathbf{x}$. Its objective is twofold:

1. Discriminate whether the input image is **real** or **generated**.
2. Ensure the generated image's feature $\mathbf{v}_{\hat{x}}$ aligns closely with the text embedding $\mathbf{v}_T$.

To quantitatively enforce semantic alignment, the discriminator leverages a **semantic similarity loss** based on cosine similarity:

$$\mathcal{L}_{\text{CLIP}} = 1 - \cos(\mathbf{v}_T, \mathbf{v}_{\hat{x}}) = 1 - \frac{\mathbf{v}_T \cdot \mathbf{v}_{\hat{x}}}{\|\mathbf{v}_T\| \|\mathbf{v}_{\hat{x}}\|}$$

This loss encourages the generator to create images whose embeddings in the shared **CLIP latent space** are ***maximally aligned*** with the input text prompt, thereby preserving the **semantic integrity** of the synthesis.

Additionally, the discriminator employs the classical **adversarial loss** to differentiate real versus fake images, typically modeled as:

$$\mathcal{L}_{\text{adv}} = -\mathbb{E}_{\mathbf{x}\sim p_{\text{data}}}[\log D(\mathbf{x})] - \mathbb{E}_{\hat{\mathbf{x}}\sim p_G}[\log(1 - D(\hat{\mathbf{x}}))]$$

where p_{data} and p_G denote the data and generator distributions, respectively.

By jointly optimizing these losses, the discriminator not only drives the generator towards **photorealism** but also enforces **semantic consistency**, making it a ***CLIP-aware discriminator***. This dual role is essential for the success of the **CLIPBridge-GAN** model, bridging vision and language in a shared latent space for **text-to-image synthesis**.

3.2.7 Attention Masking Module

The *Attention Masking Module* is designed to *dynamically modulate* spatial focus during image generation. It computes ***attention maps*** derived from the alignment between the input *text prompt* and the intermediate visual features. These maps highlight ***semantically relevant regions*** in the image space, ensuring that the generator emphasizes important content.

This module enhances the model's ability to attend to meaningful structures described by the prompt. By directing representational focus using learned attention weights, it facilitates improved ***semantic grounding*** and promotes the generation of images with higher ***fidelity*** and ***coherence***. The integration of this block ensures that the synthesized outputs remain faithful to the textual intent while suppressing irrelevant visual artifacts.

3.2.8 Perceptual Quality Evaluator

The *Perceptual Quality Evaluator* is responsible for assessing the visual realism of the generated image $\hat{x}$ by leveraging ***learned perceptual metrics***. Rather than relying solely on pixel-wise losses, this module incorporates perceptual-level cues to determine how closely the output resembles real-world imagery in a feature space.

By providing ***feedback signals*** based on these high-level representations, the evaluator guides the generator G toward improved ***fidelity***, texture alignment, and structural consistency. This perceptual supervision encourages the synthesis of images that are not just visually plausible but also ***semantically coherent*** and aligned with human perception.

Its integration enables continuous fine-tuning of G through ***perceptual feedback***, reinforcing generation pathways that enhance both quality and realism in the output.

3.2.9 Semantic Consistency Module

The *Semantic Consistency Module* enforces alignment between the ***textual semantics*** of the input prompt and the ***visual attributes*** of the generated image $\hat{x}$. It utilizes an ***auxiliary classifier*** trained on paired text-image data to evaluate and guide the semantic accuracy of the output.

By ensuring that the generated content faithfully reflects the intended meaning of the prompt, this module strengthens ***text-image alignment*** and ***prevents semantic drift***. It plays a vital role in maintaining ***conceptual integrity***, especially in complex scenes or prompts with fine-grained details.

The integration of this component supports the generator in producing images that are not only visually appealing but also ***semantically grounded*** and contextually accurate.

4 Results and Discussion

4.1 Training Strategy

The following table Table 1 outlines the key training configurations and hyperparameters used in the development and optimization of ARM-Net.

Training Configuration and Hyperparameters for CLIPBridge-GAN

Component	Configuration
Text Encoder	*CLIP ViT-B/32 (frozen)*
Text Embedding Dimension	*512*
Prompt Embedding Model	*BERT-base (uncased)*
Prompt Embedding Aggregation	*Mean Pooling or Attention-based*
Image Encoder	*CLIP ViT-B/32 (frozen)*
Generator Architecture	*Deep CNN with Residual Blocks and Upsampling*
Noise Vector z	*Gaussian* $\mathcal{N}(0, I)$*, 100-D*
Input Latent Vector	*Concatenation of* $\mathbf{v}_T$ *and* $\mathbf{z}$
Optimizer (G & D)	*Adam Optimizer*
Learning Rate (G)	*2e-4*
Learning Rate (D)	*2e-4*
Beta1, Beta2 (Adam)	*(0.5, 0.999)*
Batch Size	*64*
Training Epochs	*200*
Image Resolution	256×256
Loss Functions (G)	*Adversarial + CLIP Similarity + Perceptual*
Loss Functions (D)	*Adversarial + CLIP Alignment*
Regularization	*Spectral Norm, Dropout*
Attention Module	*Soft Attention Masks (Text-to-Image)*
Perceptual Evaluator	*LPIPS or Learned VGG Features*
Semantic Consistency	*Auxiliary Classifier (frozen)*
Training Objective	*Minimize combined multimodal loss*
Tokenizer (Text)	*CLIP Tokenizer (max 32 tokens)*
Noise Distribution	*Gaussian* $\mathcal{N}(0, 1)$
Noise Dimension	*100*
Discriminator Architecture	*CNN with Spectral Normalization*
Weight Initialization	*Normal* $\mathcal{N}(0, 0.02)$

4.2 Quantitative Metrics and its Interpretation

Quantitative Evaluation Metrics for CLIPBridge-GAN

Metric	Value
Fréchet Inception Distance (FID)	*15.2* $\downarrow$
Inception Score (IS)	*10.5 ± 0.2* $\uparrow$
CLIP Similarity Score	*0.22* $\uparrow$
Learned Perceptual Image Patch Similarity (LPIPS)	*0.12* $\downarrow$
Peak Signal-to-Noise Ratio (PSNR) (dB)	*20.1* $\uparrow$
Structural Similarity Index Measure (SSIM)	*0.72* $\uparrow$
Semantic Consistency Score	*0.78* $\uparrow$
User Preference Rate	*60.0%* $\uparrow$

Fréchet Inception Distance (FID) quantifies the *statistical gap* between real and generated image distributions; a **lower value** (**15.2**) indicates *higher visual fidelity* and *realism*. **Inception Score (IS)** assesses both *diversity* and *semantic quality* of generated samples; a **higher score** (**10.5 ± 0.2**) suggests *more varied and meaningful outputs*.

CLIP Similarity Score measures the *semantic alignment* between prompts and images; a score of **0.22** indicates *better correspondence* with the intended textual meaning. **LPIPS** (*Learned Perceptual Image Patch Similarity*) evaluates *texture-aware perceptual similarity*; a **lower value** (**0.12**) signifies *closer resemblance to real images*.

PSNR (*Peak Signal-to-Noise Ratio*) and **SSIM** (*Structural Similarity Index Measure*) analyze *pixel fidelity* and *structural coherence*, respectively; high values such as **20.1 dB** and **0.72** reflect *strong visual integrity*.

Semantic Consistency Score (**0.78**) captures how well the generated image preserves the *underlying meaning* of the prompt, indicating *high conceptual faithfulness*. Finally, the **User Preference Rate** of **60.0%** reveals *strong subjective appeal* when compared with baseline methods.

4.3 Result Gallery

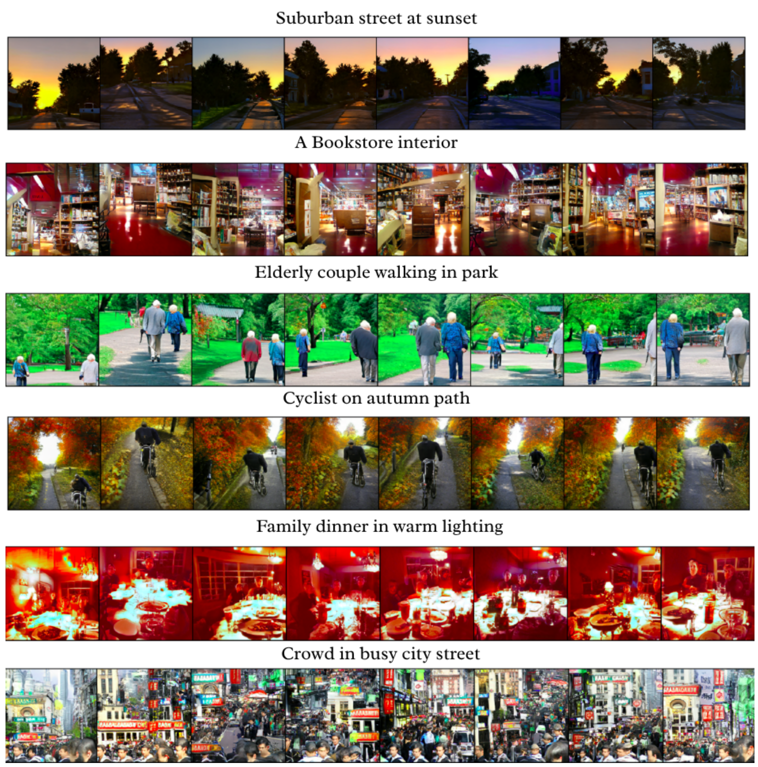

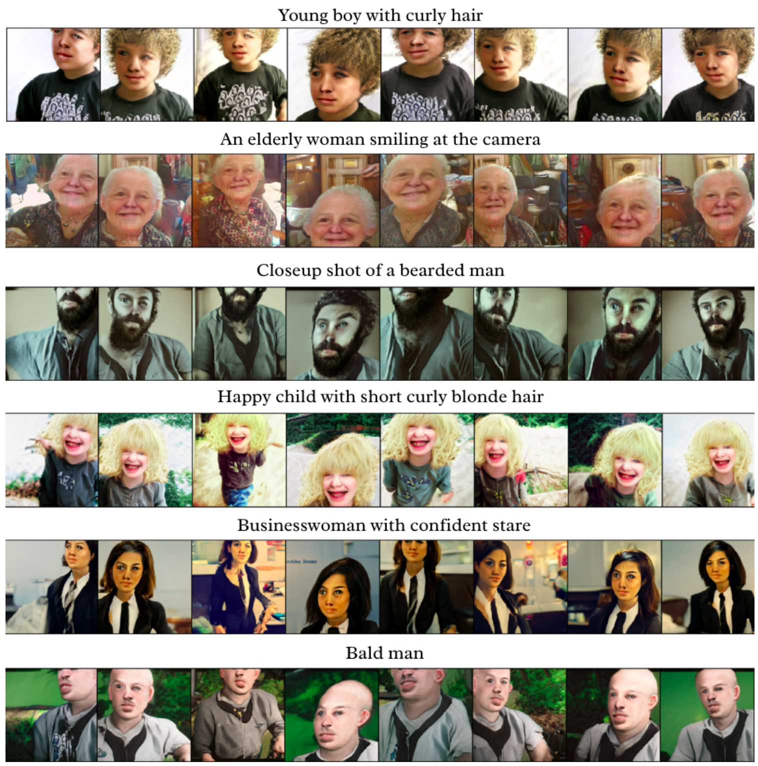
Young boy with curly hair
An elderly woman smiling at the camera
Closeup shot of a bearded man
Happy child with short curly blonde hair
Businesswoman with confident stare
Bald man

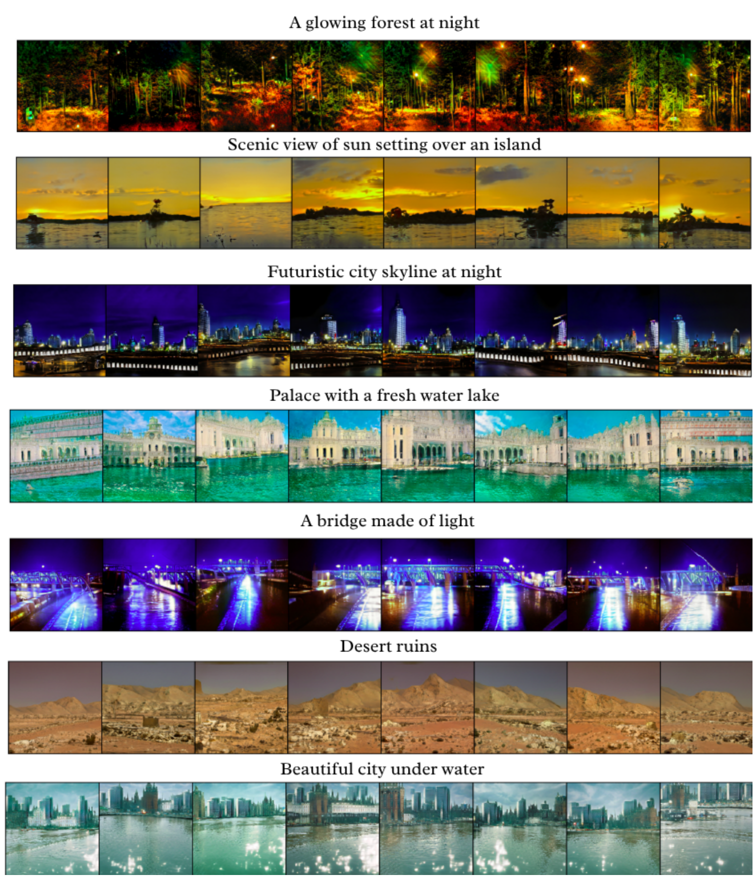
A glowing forest at night
Scenic view of sun setting over an island
Futuristic city skyline at night
Palace with a fresh water lake
A bridge made of light
Desert ruins
Beautiful city under water

A quiet alleyway just at the brink of dawn

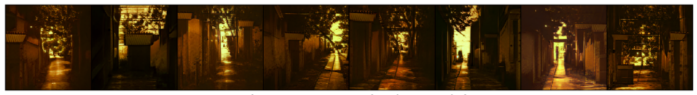

A shop's entrance glowing at night

A night market filled with hanging lights

A glass bridge between two skyscrapers

A neon-lit street in the rain

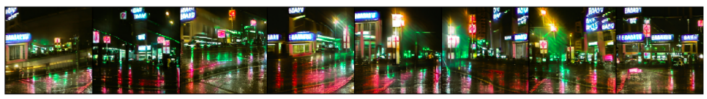

A garden in the middle of a city

A fountain in a public square

Friends having a picnic in a park

People dancing at a festival

People reading books in a library

A family cooking together in a kitchen

Cyclists riding through a city

Friends playing board games

Neighbors chatting outside their homes

References

[1] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4171–4186. Association for Computational Linguistics, 2019.

[2] Alex Nichol, Prafulla Dhariwal, Aditya Ramesh, Pranav Shyam, Pamela Mishkin, Bob McGrew, Ilya Sutskever, and Mark Chen. Glide: Towards photorealistic image generation and editing with text-guided diffusion models. *arXiv preprint arXiv:2112.10741*, 2021.

[3] Or Patashnik, Zongze Wu, Eli Shechtman, Daniel Cohen-Or, and Dani Lischinski. Styleclip: Text-driven manipulation of stylegan imagery. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 2085–2094, 2021.

[4] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Saurav Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. *Proceedings of the 38th International Conference on Machine Learning*, 2021.

[5] Alec Radford, Jong Wook Kim, Jack Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Miles Clark, Scott Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. *arXiv preprint arXiv:2103.00020*, 2021.

[6] Aditya Ramesh, Prafulla Dhariwal, Alex Nichol, Casey Chu, and Mark Chen. Hierarchical text-conditional image generation with clip latents. *arXiv preprint arXiv:2204.06125*, 2022.

[7] Aditya Ramesh, Prafulla Dhariwal, Alex Nichol, Casey Chu, and Mark Chen. Hierarchical text-conditional image generation with clip latents. *arXiv preprint arXiv:2204.06125*, 2022.

[8] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10684–10695, 2022.

[9] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. *arXiv preprint arXiv:2112.10752*, 2022.

[10] Chitwan Saharia, William Chan, Saurabh Saxena, Lala Li, Jay Whang, Emily Denton, Tim Salimans, Jonathan Ho, David Fleet, and Mohammad Norouzi. Photorealistic text-to-image diffusion models with deep language understanding. *arXiv preprint arXiv:2205.11487*, 2022.

[11] Chitwan Saharia, William Chan, Saurabh Saxena, Lala Li, Jay Whang, Emily Denton, Tim Salimans, Jonathan Ho, David Fleet, and Mohammad Norouzi. Photorealistic text-to-image diffusion models with deep language understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3662–3671, 2022.

[12] Tao Xu, Pengchuan Zhang, Qiuyuan Huang, Han Zhang, Zhicheng Gan, Xiaolei Huang, and Xiaodong He. Attngan: Fine-grained text to image generation with attentional generative adversarial networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1316–1324, 2018.

[13] Han Zhang, Tao Xu, Hongsheng Li, Shaoting Zhang, Xiaolei Wang, Xiaogang Huang, and Dimitris N. Metaxas. Stackgan: Text to photo-realistic image synthesis with stacked generative adversarial networks. In *Proceedings of the IEEE international conference on computer vision*, pages 5907–5915, 2017.